

Risk-Budgeted Mean-Variance Portfolios

Bernardo Freitas Paulo da Costa¹ **Raul Riva**² Rodrigo Targino¹

¹FGV, School of Applied Mathematics ([FGV EMAp](#))

²FGV, Brazilian School of Economics and Finance ([FGV EPGE](#))

May 15th, 2026

[FinMetrics Conference, Barcelona](#)

- Portfolio allocation is a central question in Finance
- Two major issues from the practitioner perspective:
 1. How to model the stochastic process driving returns?
 2. Given a stochastic process for returns, how to allocate money?

This paper is about the latter, not the former.

The cornerstone: Mean-Variance (MV)... but nothing comes for free

Mean-Variance (MV)

- Best trade-off between risk (variance) and expected returns
- Leads to concentrated portfolios
- Not very robust to moment estimation error

The cornerstone: Mean-Variance (MV)... but nothing comes for free

Mean-Variance (MV)

- Best trade-off between risk (variance) and expected returns
- Leads to concentrated portfolios
- Not very robust to moment estimation error

Risk Budgeting (RB)

- Assets control a maximum fraction of the total portfolio risk
- Explicit guard against portfolio concentration
- No control over expected returns

The cornerstone: Mean-Variance (MV)... but nothing comes for free

Mean-Variance (MV)

- Best trade-off between risk (variance) and expected returns
- Leads to concentrated portfolios
- Not very robust to moment estimation error

Risk Budgeting (RB)

- Assets control a maximum fraction of the total portfolio risk
- Explicit guard against portfolio concentration
- No control over expected returns

This paper: Risk-Budgeted Mean-Variance (RBMV) portfolio

- RBMV nests both the RB and MV frameworks
- Makes **explicit** the tension between expected returns and risk concentration
- Measures how much concentration you need to accept to get higher returns

1. MV and RB: a quick overview
2. Our methodology: theory and simulation results
3. Empirical application: U.S. equity market using CRSP data

Empirical takeaway:

- Our methodology indeed delivers portfolios that are less concentrated than MV
- Higher expected returns than RB
- No way around it: moment estimation is still key

- We trade d assets indexed by $i = 1, \dots, d$
- Returns r_i have an expected value $\mu_{d \times 1}$ and covariance matrix $\Sigma_{d \times d}$
- An allocation $\mathbf{v} = (v_1, \dots, v_d)^\top$ is a vector of *dollars* invested in each asset
- You have $v_0 > 0$ dollars to invest
- The dollar return for an allocation $\mathbf{v}$ is given by:

$$R(\mathbf{v}) \equiv \sum_{i=1}^d v_i \cdot r_i = v_0 \left[\sum_{i=1}^d w_i \cdot r_i \right], \quad w_i \equiv \frac{v_i}{v_0} \quad (1)$$

The MV way

- You want to minimize the variance of returns $\sigma(R(\mathbf{v}))$
- But you request a minimum expected return $\mu_{\min}^{MV}$
- The (long-only) MV allocation $\mathbf{v}^{MV}$ solves:

$$\begin{aligned} \min_{\mathbf{v} \geq 0} \quad & \sigma(R(\mathbf{v})) \\ \text{s.t.} \quad & \sum_{i=1}^d v_i = v_0 \\ & \mu(R(\mathbf{v})) \geq \mu_{\min}^{MV} \cdot v_0 \end{aligned} \tag{2}$$

- The MV portfolio is given by $\mathbf{w}^{MV} \equiv \frac{1}{v_0} \cdot \mathbf{v}^{MV}$
- This is equivalent to maximizing the return given an upper bound on volatility

The MV way

- You want to minimize the variance of returns $\sigma(R(\mathbf{v}))$
- But you request a minimum expected return $\mu_{\min}^{MV}$
- The (long-only) MV allocation $\mathbf{v}^{MV}$ solves:

$$\begin{aligned} \min_{\mathbf{v} \geq 0} \quad & \sigma(R(\mathbf{v})) \\ \text{s.t.} \quad & \sum_{i=1}^d v_i = v_0 \\ & \mu(R(\mathbf{v})) \geq \mu_{\min}^{MV} \cdot v_0 \end{aligned} \tag{2}$$

- The MV portfolio is given by $\mathbf{w}^{MV} \equiv \frac{1}{v_0} \cdot \mathbf{v}^{MV}$
- This is equivalent to maximizing the return given an upper bound on volatility

Drawback: MV portfolios can be highly concentrated on a few assets!

- If we increase v_i by \$1, how much does the portfolio risk $\sigma(R(\mathbf{v}))$ change?

Definition

The *risk contribution* of asset i to the total portfolio risk $\sigma(R(\mathbf{v}))$, is given by:

$$\mathcal{RC}_i(\mathbf{v}) \equiv v_i \cdot \frac{\partial \sigma(R(\mathbf{v}))}{\partial v_i}$$

Risk Contributions

- If we increase v_i by \$1, how much does the portfolio risk $\sigma(R(\mathbf{v}))$ change?

Definition

The *risk contribution* of asset i to the total portfolio risk $\sigma(R(\mathbf{v}))$, is given by:

$$\mathcal{RC}_i(\mathbf{v}) \equiv v_i \cdot \frac{\partial \sigma(R(\mathbf{v}))}{\partial v_i}$$

- $\mathcal{RC}_i(\mathbf{v})$ can be high for two reasons: high exposure or high asset volatility

- If we increase v_i by \$1, how much does the portfolio risk $\sigma(R(\mathbf{v}))$ change?

Definition

The *risk contribution* of asset i to the total portfolio risk $\sigma(R(\mathbf{v}))$, is given by:

$$\mathcal{RC}_i(\mathbf{v}) \equiv v_i \cdot \frac{\partial \sigma(R(\mathbf{v}))}{\partial v_i}$$

- $\mathcal{RC}_i(\mathbf{v})$ can be high for two reasons: high exposure or high asset volatility
- Since $\sigma(\cdot)$ is homogeneous of degree 1, Euler's theorem implies:

$$\sigma(R(\mathbf{v})) = \sum_{i=1}^d \mathcal{RC}_i(\mathbf{v}).$$

- Risk budgeting (RB) is a framework to limit the risk contributions $\mathcal{RC}_i(\mathbf{v})$
- We require a *risk budget* $\mathbf{b} = (b_1, \dots, b_d)^\top$ for total risk and want to invest v_0

- Risk budgeting (RB) is a framework to limit the risk contributions $\mathcal{RC}_i(\mathbf{v})$
- We require a *risk budget* $\mathbf{b} = (b_1, \dots, b_d)^\top$ for total risk and want to invest v_0

Definition (The Risk Budgeting Portfolio)

The *risk budgeting* (RB) portfolio is an allocation $\mathbf{v} \geq 0$ that satisfies $v_0 = \sum_{i=1}^d v_i$ and

$$b_i \cdot \sigma(R(\mathbf{v})) = \mathcal{RC}_i(\mathbf{v}), \quad \text{for all } i = 1, \dots, d, \quad (3)$$

- Risk budgeting (RB) is a framework to limit the risk contributions $\mathcal{RC}_i(\mathbf{v})$
- We require a *risk budget* $\mathbf{b} = (b_1, \dots, b_d)^\top$ for total risk and want to invest v_0

Definition (The Risk Budgeting Portfolio)

The *risk budgeting* (RB) portfolio is an allocation $\mathbf{v} \geq 0$ that satisfies $v_0 = \sum_{i=1}^d v_i$ and

$$b_i \cdot \sigma(R(\mathbf{v})) = \mathcal{RC}_i(\mathbf{v}), \quad \text{for all } i = 1, \dots, d, \quad (3)$$

- No clue from the definition how to *compute* $\mathbf{v}$!
- The system in (3) is a non-linear system of d equations

How to find this portfolio?

Proposition

Given a risk budget $\mathbf{b}$, any optimal solution $\mathbf{v}^*$ to

$$\min_{\mathbf{v} \in \mathbb{R}_+^d} \sigma(R(\mathbf{v})), \quad \text{subject to} \quad \sum_{i=1}^d b_i \cdot \log(v_i) \geq 0 \quad (4)$$


is proportional to the exposure $\mathbf{v}$ of the RB portfolio for risk budget $\mathbf{b}$.

- This problem is strictly convex and easy to solve
- The FOC's coincide with the RB conditions in (3)
- We can rescale the solution:

$$\mathbf{v}^{RB} \equiv \frac{v_0}{\sum_i v_i^*} \mathbf{v}^*$$

Calibrated Example

- Let $d = 5$, and $\mathbf{b} = (0.2, 0.2, 0.2, 0.2, 0.2)$ – which denotes *risk parity*
- With population moments, we compute several portfolios:




- If all you care is Sharpe Ratio, stick to MV
- Risk parity $\neq$ Equal weights
- You almost always have to visit the interior to get risk budgeting

► Calibration


Calibrated Example: Concentration

- Given a vector $\mathbf{x} = (x_1, \dots, x_n)$, we have $\text{Gini}(\mathbf{x}) \equiv \frac{\sum_{i=1}^n \sum_{j=1}^n |x_i - x_j|}{2n \sum_{i=1}^n x_i}$.
- $0 \leq \text{Gini}(\mathbf{x}) \leq 1$, and $\text{Gini}(\mathbf{x}) = 0$ if and only if $\mathbf{x}$ is uniform
- Brazilian (wealth) Gini: 0.52; South African Gini: 0.63; Swedish Gini: 0.29

(a) Portfolio weight Gini indices




(b) Risk Contributions




Calibrated Example: Adding Estimation Error

- Simulate 1 year of daily returns, estimate moments, and compute portfolios
- Plot expected returns and vols using the *population moments*

(a) Simulated portfolios in the (σ, μ) -plane



(b) Distribution of realized Gini indices for $\mathbf{w}_i$



Definition (The Risk-Budgeted Mean-Variance Portfolio)

Given a risk budget $\mathbf{b}$, an endowment v_0 , a minimum required expected return $\mu_{\min}$, and a maximum volatility bound $\sigma_{\max}$, the Risk Budgeted Mean-Variance Portfolio (RBMV) is given by $\mathbf{v} = \frac{v_0}{\sum_{i=1}^d v_i^*} \cdot \mathbf{v}^*$, where $\mathbf{v}^*$ is the solution of:

$$\begin{aligned} \min_{\mathbf{v} \in \mathbb{R}_+^d} \quad & \sigma(R(\mathbf{v})) \\ \text{s.t.} \quad & \sum_{i=1}^d b_i \log(v_i) \geq 0 & [\lambda_v] \\ & \mu(R(\mathbf{v})) \geq \mu_{\min} \sum_{i=1}^d v_i & [\lambda_\mu] \\ & \sigma(R(\mathbf{v})) \leq \sigma_{\max} \sum_{i=1}^d v_i, & [\lambda_\sigma] \end{aligned} \tag{5}$$

The corresponding portfolio weights are given by $\mathbf{w} = \frac{1}{v_0} \mathbf{v}$.

We adapt this from [Maillard, Roncalli, and Teïletche \(2010\)](#) and [Roncalli \(2014\)](#).


Our Methodology: The Lagrangian View

$$\frac{\partial \mathcal{L}(\mathbf{v}; \lambda_v, \lambda_\mu, \lambda_\sigma)}{\partial v_i} = \frac{\partial \sigma(R(\mathbf{v}))}{\partial v_i} - \lambda_v \frac{b_i}{v_i} - \lambda_\mu \left(\mu_i - \underbrace{\mu_{\min}}_{\text{moves around}} \right) - \lambda_\sigma \left(\underbrace{\sigma_{\max}}_{=0.1} - \frac{\partial \sigma(R(\mathbf{v}))}{\partial v_i} \right)$$


Our Methodology: The Lagrangian View

$$\frac{\partial \mathcal{L}(\mathbf{v}; \lambda_v, \lambda_\mu, \lambda_\sigma)}{\partial v_i} = \frac{\partial \sigma(R(\mathbf{v}))}{\partial v_i} - \lambda_v \frac{b_i}{v_i} - \lambda_\mu \left(\mu_i - \underbrace{\mu_{\min}}_{\text{moves around}} \right) - \lambda_\sigma \left(\underbrace{\sigma_{\max}}_{=0.1} - \frac{\partial \sigma(R(\mathbf{v}))}{\partial v_i} \right)$$

Connecting Risk Parity and Mean-Variance Portfolios



Divergence from Risk Budget $\mathbf{b}$



Empirical Application


- Every month, we build a long-only portfolio with the 50 largest stocks in the U.S.
- Sample: 1990-2023
- Estimate population moments $\hat{\mu}$ and $\hat{\Sigma}$ with 2 years of historical data (CRSP)
- Solve for the minimum vol portfolio and set $\sigma_{max} = \min\{\sigma_{MinVol} + 0.02, 0.2\}$
- Find the MV portfolio with the highest returns, given σ_{max}
- Set two possible values for μ_{min} :

$$\mu_{min, \text{ conservative}} \equiv \min\{\mu_{MV} - 0.05, 0.1\}$$


$$\mu_{min, \text{ greedy}} \equiv \min\{\mu_{MV} - 0.05, 0.2\}$$

- Hold 4 portfolios for 1 year: Risk Parity, $MV(\sigma_{max})$, $RBMV(\mu_{min, \text{ conservative}}; \sigma_{max})$, and $RBMV(\mu_{min, \text{ greedy}}; \sigma_{max})$

Realized Gini Index



Annualized Standard Deviation of Daily Returns



Average results

Full Sample (1990-2022)

Portfolio	Return (%)	Volatility (%)	Sharpe Ratio	Gini Index (w_i)
Risk Parity	9.68	15.69	0.84	0.16
RBMV ($\mu_{min} = 0.1, \sigma_{max} = 0.2$)	10.29	15.98	0.86	0.29
RBMV ($\mu_{min} = 0.2, \sigma_{max} = 0.2$)	10.50	16.70	0.81	0.43
Mean-Variance ($\sigma_{max} = 0.2$)	10.20	22.28	0.58	0.96

- We were able to tilt the Risk Parity portfolio towards higher returns
- Just a bit more of concentration led to higher returns and slightly higher SR
- The MV portfolio was invested in just a few assets!

► Alternative Samples

► Realized Sharpe Ratio

► Realized Returns

► σ_{max}

► μ^{MV}

Wrap-Up:

- MV is the cornerstone of portfolio allocation, but leads to high concentration
- RB tackles concentration at the expense of expected returns
- RBMV nests both, and allows you to control this trade-off in a transparent way

Wrap-Up:

- MV is the cornerstone of portfolio allocation, but leads to high concentration
- RB tackles concentration at the expense of expected returns
- RBMV nests both, and allows you to control this trade-off in a transparent way

Going forward:

- Some extensions: short-selling, transaction costs, ...
- What kind of other markets can benefit from this methodology? What examples are interesting?


Appendix and References

We assume $d = 5$ and use the following population moments for the mean returns μ , the individual standard deviations s and correlation matrix C :

$$\mu = \begin{bmatrix} 0.5 \\ 0.12 \\ 0.09 \\ 0.05 \\ 0.15 \end{bmatrix}, \quad s = \begin{bmatrix} 0.10 \\ 0.20 \\ 0.15 \\ 0.08 \\ 0.13 \end{bmatrix}, \quad C = \begin{bmatrix} 1 & 0.20 & 0.40 & 0.25 & 0.50 \\ 0.20 & 1 & -0.20 & 0.40 & 0.6 \\ 0.40 & -0.20 & 1 & -0.10 & 0.30 \\ 0.25 & 0.40 & -0.10 & 1 & 0.30 \\ 0.50 & 0.60 & 0.30 & 0.30 & 1 \end{bmatrix} \quad (6)$$


We further define $\Sigma \equiv s^T C s$.


Realized Sharpe Ratio

[► Back to Average Results](#)


-- Mean-Variance ($\sigma_{max} = 0.2$) — RBMV ($\mu_{min} = 0.1$) — RBMV ($\mu_{min} = 0.2$) -- Risk Parity

Realized Returns

[► Back to Average Results](#)



Expected Return From The MV Approach

[▶ Back to Average Results](#)

Panel B: Before the Great Financial Crisis (1990-2006)

Portfolio	Return (%)	Volatility (%)	Sharpe Ratio	Gini Index (w_i)
Risk Parity	10.79	14.42	0.89	0.15
RBMV ($\mu_{min} = 0.1, \sigma_{max} = 0.2$)	10.90	14.45	0.91	0.26
RBMV ($\mu_{min} = 0.2, \sigma_{max} = 0.2$)	10.74	14.99	0.85	0.39
Mean-Variance ($\sigma_{max} = 0.2$)	9.92	20.84	0.50	0.95

Panel C: After 2020 (2021-2022)

Portfolio	Return (%)	Volatility (%)	Sharpe Ratio	Gini Index (w_i)
Risk Parity	0.07	15.65	0.21	0.16
RBMV ($\mu_{min} = 0.1, \sigma_{max} = 0.2$)	0.93	16.00	0.27	0.46
RBMV ($\mu_{min} = 0.2, \sigma_{max} = 0.2$)	1.52	16.01	0.30	0.47
Mean-Variance ($\sigma_{max} = 0.2$)	12.62	16.31	0.84	0.96

References

-  Maillard, Sébastien, Thierry Roncalli, and Jérôme Teïletche (2010). “The properties of equally-weighted risk contributions portfolios”. In: *The Journal of Portfolio Management*.
-  Roncalli, Thierry (2014). “Introducing expected returns into risk parity portfolios: A new framework for asset allocation”. In: *Available at SSRN 2321309*.